

ML Assignment 1

October 12, 2023

1 LP3 Group B Assignment 1

1.1 Predict the price of the Uber ride from a given pickup point to the agreed drop-off location.

1.1.1 Perform following tasks:

1. Pre-process the dataset.
2. Identify outliers.
3. Check the correlation.
4. Implement linear regression and random forest regression models.
5. Evaluate the models and compare their respective scores like R2, RMSE, etc.

Dataset link: <https://www.kaggle.com/datasets/yasserh/uber-fares-dataset>

```
[1]: #Importing required libraries
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn import metrics
```

```
-----
ModuleNotFoundError                                Traceback (most recent call last)
Cell In [1], line 2
      1 #Importing required libraries
----> 2 import matplotlib.pyplot as plt
      3 import numpy as np
      4 import pandas as pd

ModuleNotFoundError: No module named 'matplotlib'
```

```
[ ]: df = pd.read_csv('uber.csv') #Reading CSV file
df.head()
```

```
[ ]: Unnamed: 0      key  fare_amount  \
0    24238194    2015-05-07 19:52:06.0000003    7.5
1    27835199    2009-07-17 20:04:56.0000002    7.7
2    44984355    2009-08-24 21:45:00.00000061   12.9
3    25894730    2009-06-26 08:22:21.0000001    5.3
4    17610152    2014-08-28 17:47:00.000000188   16.0

      pickup_datetime  pickup_longitude  pickup_latitude  \
0  2015-05-07 19:52:06 UTC      -73.999817    40.738354
1  2009-07-17 20:04:56 UTC      -73.994355    40.728225
2  2009-08-24 21:45:00 UTC      -74.005043    40.740770
3  2009-06-26 08:22:21 UTC      -73.976124    40.790844
4  2014-08-28 17:47:00 UTC      -73.925023    40.744085

      dropoff_longitude  dropoff_latitude  passenger_count
0      -73.999512      40.723217            1
1      -73.994710      40.750325            1
2      -73.962565      40.772647            1
3      -73.965316      40.803349            3
4      -73.973082      40.761247            5
```

```
[ ]: df.dtypes #Checking Datatypes.
```

```
[ ]: Unnamed: 0      int64
key      object
fare_amount      float64
pickup_datetime      object
pickup_longitude      float64
pickup_latitude      float64
dropoff_longitude      float64
dropoff_latitude      float64
passenger_count      int64
dtype: object
```

```
[ ]: df.isnull().sum() #Checking for Null Values.
```

```
[ ]: Unnamed: 0      0
key      0
fare_amount      0
pickup_datetime      0
pickup_longitude      0
pickup_latitude      0
dropoff_longitude      1
dropoff_latitude      1
passenger_count      0
dtype: int64
```

```
[ ]: df.drop(['Unnamed: 0', 'key'], axis=1, inplace=True) #Dropping first coloumn as it
    ↳is irrelevant.
df.dropna(axis=0, inplace=True) #Dropping the rows with null values.
df.head()
```

```
[ ]:      fare_amount      pickup_datetime  pickup_longitude  pickup_latitude \
0          7.5  2015-05-07 19:52:06 UTC      -73.999817      40.738354
1          7.7  2009-07-17 20:04:56 UTC      -73.994355      40.728225
2         12.9  2009-08-24 21:45:00 UTC      -74.005043      40.740770
3          5.3  2009-06-26 08:22:21 UTC      -73.976124      40.790844
4         16.0  2014-08-28 17:47:00 UTC      -73.925023      40.744085

      dropoff_longitude  dropoff_latitude  passenger_count
0          -73.999512      40.723217              1
1          -73.994710      40.750325              1
2          -73.962565      40.772647              1
3          -73.965316      40.803349              3
4          -73.973082      40.761247              5
```

```
[ ]: def haversine(lon_1, lon_2, lat_1, lat_2): #Function to find the distance using
    ↳the coordinates
        lon_1, lon_2, lat_1, lat_2 = map(np.radians, [lon_1, lon_2, lat_1, lat_2])
    ↳#Converting Degrees to Radians
        diff_lon = lon_2 - lon_1
        diff_lat = lat_2 - lat_1
        distance = 2 * 6371 * np.arcsin(np.sqrt(np.sin(diff_lat/2.0)**2+np.
    ↳cos(lat_1)*np.cos(lat_2)*np.sin(diff_lon/2.0)**2)) #Calculationg the Distance
    ↳using Haversine Formula
        return distance

df['Distance']=
    ↳haversine(df['pickup_longitude'],df['dropoff_longitude'],df['pickup_latitude'],df['dropoff_la
df['Distance'] = df['Distance'].astype(float).round(2) #Rounding-off to 2
    ↳decimals
df.head()
```


```
[ ]:      fare_amount      pickup_datetime  pickup_longitude  pickup_latitude \
0          7.5  2015-05-07 19:52:06 UTC      -73.999817      40.738354
1          7.7  2009-07-17 20:04:56 UTC      -73.994355      40.728225
2         12.9  2009-08-24 21:45:00 UTC      -74.005043      40.740770
3          5.3  2009-06-26 08:22:21 UTC      -73.976124      40.790844
4         16.0  2014-08-28 17:47:00 UTC      -73.925023      40.744085

      dropoff_longitude  dropoff_latitude  passenger_count  Distance
0          -73.999512      40.723217              1          1.68
1          -73.994710      40.750325              1          2.46
2          -73.962565      40.772647              1          5.04
```

3	-73.965316	40.803349	3	1.66
4	-73.973082	40.761247	5	4.48

```
[ ]: #Plotting a scatter plot to check for outliers.
plt.scatter(df['Distance'], df['fare_amount'])
plt.xlabel("Distance")
plt.ylabel("fare_amount")
```


```
[ ]: Text(0, 0.5, 'fare_amount')
```



```
[ ]: #Dealing with Outliers via removing rows with too large Distance and 0 or lesser
↳ distance.
df.drop(df[df['Distance']>60].index, inplace=True)
df.drop(df[df['Distance']==0].index, inplace=True)
df.drop(df[df['Distance']<0].index, inplace=True)
#Dealing with Outliers via removing rows with 0 or lesser fare amounts.
df.drop(df[df['fare_amount']==0].index, inplace=True)
df.drop(df[df['fare_amount']<0].index, inplace=True)
#Dealing with Outliers via removing rows with non-plausible fare amounts and
↳ distance travelled.
df.drop(df[df['Distance']>100].index, inplace=True)
df.drop(df[df['fare_amount']>100].index, inplace=True)
df.drop(df[(df['fare_amount']>100) & (df['Distance']<1)].index, inplace = True )
```

```
df.drop(df[(df['fare_amount']<100) & (df['Distance']>100)].index, inplace = True)
#Plotting a Scatter Plot to check for any more outliers and also to show
#correlation between Fare Amount and Distance.
plt.scatter(df['Distance'], df['fare_amount'])
plt.xlabel("Distance")
plt.ylabel("fare_amount")
```

```
[ ]: Text(0, 0.5, 'fare_amount')
```



```
[ ]: #Preprocessing the Data Using Standard Scaler in range of -1 to 1
x = df['Distance'].values.reshape(-1, 1)      #Independent Variable
y = df['fare_amount'].values.reshape(-1, 1)    #Dependent Variable
std = StandardScaler()
Y = std.fit_transform(y)
X = std.fit_transform(x)
#Splitting the data into training and testing set
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2,
#random_state=1)

[ ]: def apply_model(model): #Model to print the metrics of the various prediction
#models
    model.fit(X_train,Y_train)
    print("Training score = ",model.score(X_train,Y_train))
```

```

print("Testing score = ",model.score(X_test,Y_test))
print("Accuracy = ",model.score(X_test,Y_test))
Y_pred = model.predict(X_test)
print("Predicted values:\n",Y_pred)
print("Mean Absolute Error =", metrics.mean_absolute_error(Y_test, Y_pred))
print("Mean Squared Error =", metrics.mean_squared_error(Y_test, Y_pred))
print("Root Mean Squared Error =", np.sqrt(metrics.
→mean_squared_error(Y_test, Y_pred)))

```

```

[ ]: lr = LinearRegression()
      apply_model(lr)

```

```

Training score = 0.8023890708990102
Testing score = 0.800134921063358
Accuracy = 0.800134921063358
Predicted values:
[[-0.0856421 ]
 [ 1.40250073]
 [ 0.1072653 ]
 ...
 [-0.17833787]
 [-0.42636167]
 [-0.37124527]]
Mean Absolute Error = 0.243543639885431
Mean Squared Error = 0.19732734085539588
Root Mean Squared Error = 0.44421542167668593

```

```

[ ]: #Random Forest Model
      rf = RandomForestRegressor(n_estimators=100, random_state=10)
      apply_model(rf)

```

```

C:\Users\candr\AppData\Local\Temp\ipykernel_7216\3813684645.py:2:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples,), for example using
ravel().

```

```

      model.fit(X_train,Y_train)

Training score = 0.8250567049453948
Testing score = 0.7931312012692804
Accuracy = 0.7931312012692804
Predicted values:
[-0.10304075  1.80284551  0.08764113 ... -0.21391608 -0.42011423
 -0.37785255]
Mean Absolute Error = 0.24703500001737674
Mean Squared Error = 0.20424213262599705
Root Mean Squared Error = 0.4519315574575392

```

```

[ ]:

```